



Cut the crap: a critical response to “ChatGPT is bullshit”

David Gunkel^{1,2} · Simon Coghlan³

Published online: 18 April 2025
© The Author(s) 2025

Abstract

In a recent thought-provoking essay called “ChatGPT is Bullshit,” Hicks, Humphries and Slater call such large language models (LLMs) “bullshitters” and “bullshit machines.” Unlike the term “bullshit,” they argue, commonly used anthropomorphic terms such as “hallucination” and “confabulation” mispresent LLMs and sow confusion that could be socially harmful. This paper criticizes their essay in two steps. First, its reliance on Harry Frankfurt’s classic characterization of bullshit as indifference to truth, though understandable and compelling in one sense, risks misrepresenting LLMs. Second, the argument is too quick to jettison anthropomorphic terms like hallucination and confabulation, which might prove useful metaphors for understanding generative AI. Exploring language to articulate good ways of understanding LLMs is indeed a socially important task, one benefitting from critical open-mindedness, some historical awareness, and a nuanced approach to how various words used to describe AI can operate. This paper attempts to contribute to this task by questioning the wisdom of categorically calling bullshit on ChatGPT.

Keywords Bullshit · ChatGPT · Confabulation · Frankfurt · Generative artificial intelligence · Hallucination · Large language models · LLM

Introduction

Large language models (LLM), such as OpenAI’s ChatGPT, are notorious for producing output that sounds good but is not true. These problems have been called *hallucinations* or even *confabulations* due to the fact that the generated output looks legit—it is grammatically correct and often sounds convincing or appears to be plausible—but is entirely fabricated, fictional, or false. But the terms “hallucination” and “confabulation” have not sat well with critics, who find these and similar anthropomorphic concepts to be fraught with problems (Emsley, 2023; Edwards, 2023; Smith et al., 2023). In response to this, Michael Townsen Hicks, James Humphries, and Joe Slater have now offered us an

innovative essay that advances what they argue is a better characterization—*bullshit*. In fact, the authors not only call LLMs “bullshit machines,” they “call bullshit on ChatGPT” (Hicks et al., 2024, p. 2; see also Slater et al. 2024).

The essay has been undeniably popular. In a short time, it amassed over 800,000 views and 54 citations, and news of its publication spread quickly across social media with numerous posts and reposts. Leveraging Harry Frankfurt’s little book on the subject, the authors set out to demonstrate that bullshit is a better characterization than either hallucination or confabulation and, more importantly, that LLMs in general and ChatGPT in particular should be characterized as bullshitters. But, as we will argue in this critical response essay, while calling bullshit on ChatGPT is understandable and compelling in one sense, it may be seriously questioned.

We will proceed in three steps or movements: (1) We begin by noting how the title “ChatGPT is bullshit” is a kind of *Zeitgeist* statement, thus demonstrating how Hicks et al. pick-up on and elaborate a thesis that had already been circulating in the networks of contemporary culture. (2) We then reexamine the concept of bullshit that is appropriated from Frankfurt and demonstrate how his formulation, which is presented in the book *On Bullshit*, can be misleading, especially when uncritically extended to LLMs. Our argument

✉ Simon Coghlan
simon.coghlan@unimelb.edu.au

David Gunkel
dgunkel@niu.edu

¹ Northern Illinois University, DeKalb, USA

² Uczelnia Łazarskiego, Warsaw, Poland

³ University of Melbourne, Melbourne, Australia

here employs an important distinction, introduced by one of Frankfurt's critics, between activity-centered bullshit and output-centered bullshit (Cohen, 2002). (3) Finally, we question the argument that "bullshit," "bullshitters," and "bullshit machines" are better metaphors for describing LLMs than fabrication, hallucination, or confabulation.

As Hicks et al. and other commentators have claimed (Smith et al., 2023), evaluating appropriate terminology to describe various forms of artificial intelligence (AI) is a socially important task. Critically examining how we have and should describe AI—including the use of language that makes comparisons to the abilities of humans and animals—may also improve our understanding of and relations to this strange new technology. This linguistic and conceptual endeavor is therefore worthwhile in its own right. Our aim in contributing to this exploration of language in the context of technology is not to show that Hicks et al.'s detailed exploration is totally mistaken, but instead to demonstrate that even greater nuance is required in assessing such descriptive terminology.

Bullshit literature

Applying the term "bullshit" to ChatGPT and doing so by making reference to and mobilizing Frankfurt's characterization of the concept is something that had already been circulating on the Internet and presented in other scholarly research that was available either prior to or in parallel with Hicks et al.'s article. Though Hicks et al. do not explicitly refer to or cite these other publications, we take note of them here in an effort (1) to identify how the argument put forward by Hicks et al. does not so much introduce the connection that is made between ChatGPT and Frankfurt's book *On Bullshit* but contributes to and even elaborates on an idea that had already been available and in circulation; and (2) to help us identify and focus on the unique contributions that their article adds to this formulation. This first section, then, serves as a kind of critical review of the LLM bullshit literature.

Already in December of 2022, Arvind Narayanan and Sayash Kapoor posted the following to their blog *AI Snake Oil*: "The philosopher Harry Frankfurt defined bullshit as speech that is intended to persuade without regard for the truth. By this measure, OpenAI's new chatbot ChatGPT is the greatest bullshitter ever. Large Language Models (LLMs) are trained to produce plausible text, not true statements." In these three sentences, Narayanan and Kapoor not only call ChatGPT a bullshitter but do so by making explicit and direct connections to Frankfurt's characterization of the term. The statement has been subsequently included in their

book *AI Snake Oil*, which was published in September of 2024 (Narayanan and Kapoor 2024, 139).

Also in December of 2022, Valliappa Lakshmanan posted an article to LinkedIn with the title "Why Large Language Models like ChatGPT are Bullshit Artists, and How to Use Them Effectively Anyway." Lakshmanan's title is not only similar to that of Narayanan and Kapoor's blog post but also leverages Frankfurt's book in his (Lakshmanan's) characterization of ChatGPT as a bullshitter: "What do you call a system that sounds confident about what it says regardless of whether what it's saying is true? I'll give you a hint. Harry Frankfurt wrote an entire book cautioning against a person who 'does not reject the authority of the truth, as the liar does, and oppose himself to it. He pays no attention to it at all.' He calls such a person a bullshitter."

A month later, in January 2023, Jürgen Rudolph, Samson Tan, and Shannon Tan published an article in the *Journal of Applied Learning and Teaching* titled "ChatGPT: Bullshit Spewer or the End of Traditional Assessments in Higher Education?" In this essay, Rudolph et al. not only ask whether ChatGPT is a "bullshit spewer" but anchor their use of the term in an earlier publication from 2020: "At the other end of the spectrum of responses, Marcus and Davis (2020) proclaimed that GPT-3 is "a fluent spouter of bullshit" and "not a reliable interpreter of the world."

A month later, in an article for *Vice* published in February 2023 and titled "ChatGPT is a Bullshit Generator Waging Class War," Dan McQuillan offered the following argument:

ChatGPT is, in technical terms, a "bullshit generator." If a generated sentence makes sense to you, the reader, it means the mathematical model has made a sufficiently good guess to pass your sense-making filter. The language model has no idea what it's talking about because it has no idea about anything at all. It's more of a bullshitter than the most egregious egoist you'll ever meet, producing baseless assertions with unfailing confidence because that's what it's designed to do.

Here, McQuillan (2023) not only applies the term "bullshit" to ChatGPT but provides a characterization and description that aligns with what is developed by Hicks et al.

Similar uses of the term were provided in July of that same year by Alberto Romero (2023) in a blog post for *The Algorithmic Bridge*, titled "ChatGPT: A Bullshit Tool For Bullshit Jobs," and in October with Kolassa (2023) posting "ChatGPT as a Generator of Mindless Bullshit" on LinkedIn. In this short article, Kolassa not only uses the term "bullshit" but makes a statement that, by way of the *Wikipedia* entry on the subject, again connects the dots to Frankfurt's book: "it takes only two minutes to fire up ChatGPT

and have it vomit out a bucketful of bullshit (and I am using this term in its technical meaning: https://en.wikipedia.org/wiki/On_Bullshit).”

What is important—and what we will focus on—is not the mere claim that ChatGPT is a bullshit generator but how and why calling ChatGPT a bullshit generator is presented as being an arguably better and more accurate description of LLM operations than the other terms that have also been in circulation. In other words, what makes Hicks et al.’s essay interesting and important is not that they were the first or the only researchers to see a connection between the generated content of LLM AI and Frankfurt’s concept of bullshit, since that suggestion resonates with what a lot of people were already thinking. Rather, what Hicks et al. contribute is a detailed argument for selecting this one particular term over the other ones that have been available to us. What follows, therefore, an analysis that is directed at terminological precision, truthful communication, and transparency.

On the concept of bullshit

Hicks et al. (2024) argue in favor of using the term “bullshit” to describe what ChatGPT generates and “bullshitter” and “bullshit machine” as what ChatGPT is. In doing so, they directly oppose and argue against the use of terms like “hallucination,” “lying,” and “confabulation” to name and characterize what we experience as inaccurate and false output. The existing terminology (especially, “hallucination” and “confabulation”), they protest, are *anthropomorphic* (Hicks et al., 2024, 8–9). Consequently, such terms wrongly imply that LLM’s possess human-like psychological capabilities, such as a concern with and an ability to discern truth (which we might say is a necessary condition of human activities like lying and fabrication), the capacity to “believe or perceive” reality (necessary for having hallucinations), and the faculty of “remembering” and “convey[ing] accurate information” (necessary for being able to confabulate) (Hicks et al., 2024, p. 9).

Instead, the authors point out, all that LLMs do (or are capable of doing) is to predict tokens or next words in probabilistic sequences of linguistic signifiers, which is very far from attending to the truth of what these words say or mean. And it is because of this that the existing metaphorical terms like “hallucination” are “inapt” (Hicks et al., 2024, p. 1). Indeed, they not only fail to capture what LLMs actually do (or are incapable of doing) but also mislead the public and play into the “overblown hype” (Hicks et al., 2024, p. 9) that has been associated with the marketing discourse of contemporary AI. Bullshit, by contrast, is therefore proposed as a better and more responsible way to characterize LLMs like ChatGPT.

Because Hicks et al.’s argument is squarely situated on Harry Frankfurt’s conception of bullshit, a brief recap of his characterization of this typically vulgar word will be useful prior to evaluating its application to LLMs and generative AI. In his book on the subject, Frankfurt (2005) distinguished bullshitting from lying. Whereas lying entails knowing or believing the falsehood of what is stated (Mahon, 2016), bullshitting consists in a complete indifference to a statement’s veracity. The liar believes in or knows the truth and intentionally misrepresents it; the bullshitter has no sincere regard for whether what they state is, in fact, true or not. What the bullshitter says or claims might be false or it might be true—in both cases it is bullshit.

Additionally, Frankfurt (2005, 54) characterizes a bullshitter as someone who misrepresents their agenda, which is to say that they deceive or mislead their audience into thinking that they are “endeavoring to communicate the truth.” In his view, then, bullshitting is dangerous, even more dangerous than lying (Frankfurt, 2005, p. 61) or even mere incoherence (Frankfurt, 2002, p. 343). That is because, unlike the person who routinely and intentionally misrepresents the truth, or the person who constantly makes incoherent claims while trying to say something true, the bullshitter is completely indifferent to truth.

According to Hicks et al. (2024, 1), LLMs are bullshitters, in Frankfurt’s sense of the word, “because these programs cannot themselves be concerned with truth, and because they are designed to produce text that looks truth-apt without any actual concern for truth.”¹ This basic indifference to truth they call “soft” bullshitting and the results or output that is produced “soft bullshit.” They further propose, though with less confidence, that LLMs might be regarded as “hard” bullshitters that produce “hard bullshit,” insofar as LLMs can also be plausibly said to perform the deception of misrepresenting their agenda as communicators of the truth. Thus, when ChatGPT proclaims in what appears to be authoritative, unequivocal language that such-and-such is the case, it deceives us about what it is up to. “The fact that it is designed to give the impression of concern for truth qualifies it as attempting to mislead the audience about its aims, goals, or agenda” (Hicks et al., 2024, p. 2), which is “to portray itself as a ‘normal’ interlocutor like ourselves” (Hicks et al., 2024, p. 7). Consequently, the authors claim that since LLMs necessarily engage in bullshitting—in at least the first sense of “soft” bullshit and possibly also in the second sense of “hard” bullshit—everything they produce is bullshit.

¹ One could argue for enlarging Frankfurt’s definition to cover the phenomenon of not so much being indifferent to the truth as grossly failing to show epistemic responsibility about what is said (see also Johnson, 2010). The argument might then be that LLMs lack epistemic responsibility.

By relying on this characterization, Hicks and company expose themselves to the criticisms that have been directed at Frankfurt's position (Cohen, 2002; Levy, 2023). Despite offering a disclaimer in the opening pages of *On Bullshit* that bullshit is a vast phenomenon and complex concept, Frankfurt (2005, 2–3) nonetheless implied that his own characterization is the dominant one and that his formulation captures bullshit's *essence* (Frankfurt, 2005, p. 34). But this view, as we will demonstrate, is too narrow and limiting as it overlooks and thus sidelines important, if more obvious, ways that people use the term. Indeed, Frankfurt (2002, 340) had already agreed with G.A. Cohen's (2002, 324) criticism of the original essay "On Bullshit" (which Frankfurt had published in the *Raritan Quarterly Review* in 1986) that there are in fact *two* main types of bullshit: *activity-centered* and *output-centered* (see also Levy, 2023, p. 7; Pennycook et al., 2015).²

An activity-centered conception of bullshit like Frankfurt's³ focuses on the agent's "mental state" (Frankfurt, 2002, p. 340) or their attitude regarding what they say or write. We could understand this mental state or attitude in different ways. Thus, a broad activity-centered conception might include the agent's lack of concern, interest, diligence, or care regarding the truth of their statements. In contrast, an output-centered conception of bullshit focuses on the utterances or statements themselves. Such statements, we might say, can exhibit features like factual falsity, exaggeration, illogicality, damning shallowness, and pseudo-profundity (Pennycook et al., 2015), and may range from moderate error to utter nonsense and rubbish. Importantly, output-centered bullshit can be identified without knowing the mental state or attitude of its progenitor.

Adopting an output-centered approach seems to render the claim that LLMs can produce bullshit very plausible. LLMs, as Hicks et al. (2024, 1) point out, typically produce legible and meaningful text that looks "truth-apt" rather than text that is entirely meaningless or gibberish.⁴ Yet notoriously, this content can include false and/or inaccurate claims, illogical propositions, and faulty "reasoning." According to the output-centered characterization, these would or could count as bullshit, perhaps to varying degrees.

If one wants to claim that ChatGPT bullshits, the output-centered characterization arguably offers advantages over

the activity-centered characterization adopted by Hicks et al. First, there is an active, open, and sophisticated debate (Levinstein & Herrmann, 2024; Millière & Buckner, 2024) about whether LLMs can, for example, "track truth," (Hicks et al., 2024, p. 9) even though it seems intuitively obvious that they cannot grasp or understand truth. The mere existence of this dispute already complicates Hicks et al.'s (2024, 9) assertion that LLMs can at best only track truth "incidentally" due to the fact that they are mere predictors of linguistic tokens or likely next words. It might at least be seriously argued that this activity is a part of what is involved with speaking the truth and/or that LLMs somehow really do "track truth"—and such a position, even if mistaken, requires a considered reply.

Second, and as we shall explain in this and the next section, it remains to be seen whether LLMs can, in fact, bullshit in Frankfurt's activity-centered conceptualization. Perhaps they can, but this still requires further argument. For these two reasons, adopting an output-centered approach to describe and to warn us of the perils of LLM generated content faces fewer obstacles than trying to deploy an activity-centered formulation. In other words, it might be easier and even sufficient to say that LLMs are apt to produce varying degrees of bullshit content.

However, it is clear that Hicks et al. would not be satisfied with an output-centered approach that uses the term "bullshit" only to describe forms of bad and false content produced by LLMs. For the authors insist that *everything* LLMs like ChatGPT generate is bullshit. On Frankfurt's understanding of bullshit, everything that is said without concern for whether it is true or not is necessarily and without exception bullshit regardless of its actual truth value or merit. Indeed, in his reply to Cohen's criticism of the essay version of the argument, Frankfurt (2002, 340) (as Hicks et al. themselves note) went further still, claiming that if a non-bullshitter repeats something that a bullshitter has said, then that too is bullshit, even if the person repeating it sincerely believes it to be true. And this would be the case even if what is repeated happens to be not only true but objectively comprehensive, insightful, nuanced, and/or original.

There is certainly something to be said for this kind of blunt formulation. A human bullshitter's indifference to the truth, or their phoniness, or their epistemic irresponsibility, are often objectionable features in and of themselves. Furthermore, the average bullshitter possessing such qualities is likely to produce false and worthless statements with considerable regularity, and that is a good reason for not trusting or relying on what they say. Something analogous to this last point can be said of LLMs. Calling everything LLMs produce bullshit at least draws attention to the fact that the way LLMs work is connected to their inherent ability to

² Though it must be said that in focusing on Cohen's (2002) relatively niche (output-centered) understanding of bullshit as "unclarifiable discourse," Frankfurt (2002) joined Cohen in downplaying the common and straightforward (output-centered) understanding of bullshit as untrue or exaggerated statements, of which there is no shortage.

³ To reflect this activity or agent-oriented focus, Frankfurt's *On Bullshit* might well be titled *On Bullshitting*.

⁴ Nor is ChatGPT prone to producing Cohen's (2002) "unclarifiable" bullshit.

generate false or misleading content. In that sense, the claim is understandable and even seems compelling.

Nonetheless, the above formulation may be too coarse and imprecise. In a perfectly intelligible and commonsensical way, human bullshitters—like Hicks et al.’s (2024, 7) “unprepared student or the incompetent politician”—can and do say things that are not entirely bullshit. Because, unlike the incorrigible liar, they are indifferent to the truth, these individuals might indeed inadvertently come out with truth and speak “good sense” (Cohen, 2002, p. 331) in the very process of bullshitting. Furthermore, some human bullshitters are highly sophisticated, brilliantly stitching together things others have said and even inventing compelling statements and impressively cogent arguments (Levy, 2023).⁵ If someone were to bluntly insist that such statements or arguments are just bullshit irrespective of how true and cogent they are, we might be sympathetic but still reluctant to agree. We might want to reply that even if those utterances are bullshit in one sense, they are decidedly not bullshit in another.

People use the term “bullshit” just as commonly or naturally in the output-centered sense as in the activity-centered sense. Frankfurt’s analysis captures something vital about the concept of bullshit, but it does not fully capture its “essence” or exhaust its important and basic meanings. Consequently, categorically calling bullshit on everything the bullshitter says, as is the case for those advancing a purely activity-centered characterization, can be misleading. To be sure, Frankfurt (2005, 48) for one was perfectly clear from the start that what the bullshitter comes out with “need not be false.” But our point is that the blunt and blanket labelling as “bullshit” everything the person bullshitting says or writes, regardless of its objective quality, can lead to confusion. While it might be true that the majority of human bullshitters have little to say that is of value, that is not a necessary or exceptionless rule about them.

As it turns out, this matters a good deal for how we think about and evaluate those unusual entities called LLMs. On an output-centered understanding of the term, bullshit is characterized as statements or claims that are false, exaggerated, nonsense, garbage, etc. But if they are indeed indifferent to the truth, then LLMs are equally capable of producing both what we call bullshit and statements that we would evaluate as true. And, as we have sometimes been astonished to see, LLMs, even if lacking a concern for truth, can produce written statements or sentences that are not only

“accurate” but also “comprehensive” (Hicks et al., 2024, p. 3), insightful, nuanced, and original. While there is likely a limit to what they can do in these respects (Marcus, 2024), LLMs like ChatGPT are still improving along these dimensions. Consequently, it seems seriously misleading to categorically insist that everything ChatGPT and other LLMs output is bullshit. And for this reason, it also seems misleading to call ChatGPT a “bullshitter” or a “bullshit machine” (Hicks et al., 2024, p. 7). These latter expressions tend to suggest that ChatGPT can only produce untrue or worthless outputs. And because these expressions tend to convey that impression even when that implication is not intended, they are liable to be misunderstood.

At a minimum, claims that ChatGPT or some other similar LLM is a bullshitter, bullshit machine, or producer of bullshit need qualification. But even when we are careful to say that these expressions are only meant in the activity-centered sense, there is arguably something still not quite right about calling LLMs bullshitters or bullshit machines. If they are bullshitters, LLMs (or at least the better LLMs) are bullshitters in a rather unusual sense. These LLMs would be more like a human being who, while making statements without caring about their truth value or worth, nevertheless very often state things that are true, factual, helpful, logical, backed by reasons, insightful, creative, superhumanly knowledgeable, and so on.⁶

We should emphasize that this argument is not inconsistent with being acutely wary of the fact that LLMs sometimes or even many times generate content that is not correct or true. Such falsehoods rightly affect our trust in LLMs, and this legitimate distrust accentuates worries about the tsunami of machine-generated content heading our way (Kirschenbaum, 2023). Furthermore, making the public aware of any propensity LLMs have to output questionable content can be regarded as practically important and responsible. But unequivocally categorizing everything generated by an LLM as bullshit and categorically labeling LLMs bullshitters and bullshit machines could undermine a more practical and balanced understanding of this technology and its actual social impact and consequences.

⁵ This is not to deny that the fact they are bullshitters and can seriously undermine or reduce the value of *some* of the things they say. For example, we would be unlikely to think they can impart genuine wisdom, insofar as wisdom requires, at minimum, a real concern for the truth about things that matter. LLMs and other AI cannot be wise in this sense.

⁶ In fact, it might be considerably stranger to call such an LLM a bullshitter than it is to call this (highly unusual) human being a bullshitter. For human beings are in the fullest sense able to “track truth” (Hicks et al., 2024, p. 9) and to really care about it, whereas an LLM is not. Furthermore, LLMs are trained on some factual data and then given further reinforcement training to produce accurate and worthwhile content (even if they are apt to fail). We make further use of these points in the next section.

Competing metaphors

As an integral component of their effort to extend the term bullshit from human beings to LLM AI, Hicks and colleagues oppose and argue against doing the same with other terms like “hallucination” and “confabulation.” In the previous section, we considered problems with Hicks et al.’s conception of bullshit as applied to LLMs. We now want to ask the following question: If it is after all legitimate to extend an activity-centered conception of bullshit from humans to LLMs, why is a similar move not possible for at least some of those other activity-centered terms variously deployed by AI practitioners, critics, and journalists to describe ChatGPT and other LLMs? In other words, what justifies the extension and use of one anthropocentric term like “bullshit” and not other terms like “hallucination” or “confabulation?”

To start with, it is fair to say that terms like “hallucination” and “confabulation” have their usual and primary sense in relation to humans and their psychological, cognitive, and phenomenal sensory abilities. It is arguably and equally true that LLMs lack these capacities.⁷ Hicks et al. (2024, 8–9) seem to think that these facts alone decisively demonstrate that LLMs cannot hallucinate or confabulate. For example, they object that “the term hallucination anthropomorphises the LLMs” and that “LLMs do not perceive, so they surely do not ‘mis-perceive’” (Hicks et al., 2024, p. 8). But if anthropomorphism is the objection, then the move to assign to LLMs a capacity for bullshitting seems to encounter the same or at least a substantially similar problem.

Frankfurt’s concept of bullshit is predicated on how *humans* talk and behave. His conception refers to agents who make statements without concern for their truth, but crucially these are agents who *could be concerned about truth* (cf. Johnson, 2010). We might agree with Hicks and colleagues that LLMs cannot be concerned for truth even if they inadvertently produce truthful statements. But this is a different sort of indifference to that of the human bullshitter. The human bullshitter is an agent who could be concerned about the truth of their utterances but is not. The LLM cannot be indifferent in *that* sense due to the fact that it neither can nor cannot care about truth.

The difference in meaning here may manifest in different ways. Examples in *On Bullshit*, such as the phenomena of bullshitting in advertising and politics, remind us of the fact that we often feel anger, disdain, disappointment,

annoyance, and exasperation towards the bullshitter. Those responses reveal a recognition that a human bullshitter could indeed care about the truth of their claims but for whatever reason does not exercise that care. Our reactions in those contexts depend on such agents being “conscious beings” (Frankfurt, 2005, p. 66) of the kind that can distinguish truth from falsehood, honesty from dishonesty, fake talk from straight talk.

Not all agents, even conscious ones, can bullshit. For example, we surely do not have those same reactions of disappointment and anger towards, say, a young child learning their first language who repeats the falsehoods and nonsense they pick up from television and older siblings. If such a child is bullshitting, they are bullshitting in a way that is very different from that of adults. Clearly, LLMs are in one important sense closer to the mimicking, copycat young child⁸ than to the bullshitting adult.

There is a reason that Emily Bender, Timnit Gebru, Angelina McMillan-Major, and Margaret Mitchell (2021) used the phrase “stochastic parrots” in their well-known paper on LLMs, as parrots (and other kinds of “talking birds”) apparently string together words without knowing (or not knowing) the truth of what they say. But we typically would not call these birds bullshitters.⁹ Our point here is that the extension of Frankfurt’s concept of bullshit to LLMs is not as straightforward or as self-evident as Hicks et al. seem to assume given their apparent opposition to anthropomorphic language. And this remark would apply as much to their notion of “soft” as it would to “hard” bullshit.¹⁰

Hicks et al. could reply to this critique in the following way: They could argue that despite the important differences between the “utterances” of LLMs and human beings, there is sufficient overlap between the performative activity of the two for the term bullshit to still make good sense when used to describe the operations of the LLM. Furthermore, some of the statements that are made in their essay

⁷ Even if we admit that LLMs can reason, understand, or track truth (Millière & Buckner, 2024) in some sense (e.g. a computational one), it need not follow that LLMs can grasp truth or perceive reality as humans and some animals do. This obviously opens another philosophical can of worms which we will not have the time or space to pick through in this essay.

⁸ Of course, LLMs are also very different from young children, who, unlike ChatGPT, are in the process of becoming genuine language users with a conceptual grasp of truth.

⁹ We should, however, say that it is arguably ignorance and human chauvinism that leads people to believe that “talking birds” output speech sounds without any understanding, just as LLMs do (Pepperberg, 2012). Unlike LLMs, parrots and other animals have a sophisticated mental life. Exactly how “language-producing” parrots and other animals are doing something similar to and different from humans and AI requires unbiased investigation. In fact, the use of the term “parrot” by Bender et al. is another instance of the problem of extending existing words to LLMs without first critically evaluating the full consequences of that seemingly simple linguistic operation.

¹⁰ Since “hard” bullshit requires assigning to LLMs something like an intention to misrepresent their agenda in the making of true statements, such “hard” bullshit looks harder to justify on Hicks et al.’s own assumptions and requirements about, say, avoiding erroneous anthropomorphism. But for the reasons given, “soft” bullshit also faces what we might call the “objection from anthropomorphism.”

appear to concede that bullshit might be regarded as an apt metaphor, rather than as a strictly literal description of the operations or behaviors of LLMs.¹¹ In other words, it seems entirely possible that an activity-centered form of bullshit might apply to LLMs in a related albeit different way from how it applies to humans.

As we will soon show, we are not at all against the idea that certain anthropomorphic terms, including “hallucination” and “confabulation,” might have a not-strictly-literal or even a metaphorical or analogical application to LLMs (see, Smith et al., 2023). But Hicks et al. cannot have it both ways. Either they reject all these “anthropomorphic” and “metaphorical” terms that are derived from and possess connections to human (or animal) psychological and cognitive abilities—and that would include not only “hallucination” and “confabulation” but also “bullshitting.” Or they remain open to accepting and employing these terms as more or less useful forms of description. The problem for Hicks et al. is that they seem to want to do both. They want to dismiss one set of terms (i.e. “hallucination” and “confabulation”) as anthropomorphic metaphors that are therefore inappropriate for describing the operations of LLMs and then accept the other anthropomorphic metaphorical term (“bullshitting”) as somehow insulated from or situated outside of these apparent problems.

The sensible option, then, appears to be to avoid summarily slamming the door on a range of activity-centered or agent-orientated terms that are “anthropomorphic.” Such terms need not be *problematically* anthropomorphic (Coghlan, 2024), for they are not to be seen as having precisely the same meanings when applied to machines as when applied to humans. As it happens, people already commonly apply to computer technology and AI a range of terms that originally applied only to humans and animals and that have a kind of metaphorical or extended meaning in that new context. Examples are “agent,” “autonomous,” “learning,” “memory,” “prediction,” “inference,” and “decision.” All of these and other originally human-centered terms have proven to be useful ways of describing technology, even though we might have baulked at those words in the first instance. In fact, it may be impossible to say many important things about *artificial* intelligence without using words and phrases that have previously been limited to natural intelligence seen in human beings and nonhuman animals.

As we say, we should be cautious about automatically excluding words like “hallucination” and “confabulation” simply on the grounds that they are anthropomorphic metaphors. However, Hicks et al. (2024, 9) have a possible reply to this, for they could claim that these terms are not merely

anthropomorphic metaphors but *defective* anthropomorphic metaphors.¹² We want first to agree with the thought that it is wise to try, as others too have recommended, to use “metaphoric language with precision” (Smith et al., 2023) when seeking to describe technology like AI. Not all words are created equal, and one metaphor or analogy may be better, less misleading, or more expedient than another for describing and understanding LLMs. Yet we are less certain that words like “hallucination” and “confabulation” are defective metaphors for that technology.

We have already noted that metaphors can be apt even when there are significant differences between the object of the metaphor and the thing from which the metaphor originated, as we see with terms like “machine learning.” It might also be useful to distinguish between two levels of description: (1) those descriptions that are orientated towards LLMs’ *internal processes* operating “under the hood” (Levinstein & Herrmann, 2024); and (2) those descriptions related to the *external behaviors* or generated results that users encounter and experience. Potentially, some descriptions may be dubious at one level but not the other. For example, it *might* be less plausible to say that LLMs hallucinate at the internal computational level but more plausible to say that they sometimes do produce hallucinations at the external behavioral level in the content that is generated and then encountered and experienced by users. In any case, it seems prudent to be circumspect when it comes to both word selection and rejection. We therefore conclude this section with a brief reevaluation of some of the existing terminology that had been rejected by Hicks et al. These terms are certainly as contestable as bullshit and are clearly open to further debate. But keeping them in play can provide for understanding and insight that the term “bullshit” risks foreclosing.

*Fabrication.*¹³ This description, though a metaphor drawn from human behavior, appears to make some sense in the context of an LLM that generally, or sufficiently often, provides sensible and accurate output (even if there is no deeper process of fact-checking occurring in the model). When, for example, an LLM gives a false or inaccurate response in the context of an otherwise factually accurate conversation, we might call that “fabrication.” And the fact that the LLM may subsequently identify its own mistake (Zhang et al. 2023) and correct it (and even apologize for generating the falsehood) may help create the sense that it was prone to or guilty of fabricating the false statement in the first place. There are also instances where fabrication is both deliberate and desired, as is the case when users prompt LLMs to generate various forms of fiction. Fiction

¹¹ For example, the authors’ claim that hallucination “is the wrong metaphor” (Hicks et al., 2024, p. 9) invites the impression that bullshit might be the right metaphor.

¹² Thanks to a reviewer for highlighting this point.

¹³ While this term is not used in Hicks et al. (2024), it is another anthropomorphic term that some are likely to object to.

is clearly not reducible to nor is it the same as false or misleading information. There is, as demonstrated in the work of 20th century literary theory, a truth in fiction that is not reducible to mere factual reporting (Ricoeur 1984, Derrida, 1992). Thus, unlike “bullshit,” which by comparison risks excluding any and all forms of fictional content, “fabrication” validates this important aspect of human communication and truth making.

Hallucination. Unlike “fabrication,” the term “hallucination” connotes not just any false output but wildly unexpected kinds of content that in some instances might even be called “creative” or “inventive.” The term “hallucination” has not only been explored by psychologists and philosophers (Allen, 2014), but it actually has a storied history in the field of AI and artificial neural networks (ANNs) in particular. In the early 2000s the word was commonly used (positively) to describe beneficial features of computer vision systems, specifically super-resolution, image inpainting, and image synthesis (Maleki et al., 2024). In these cases, the advantage of a term like “hallucination” is that it takes things to a level beyond error or even bullshit and into the realm of the fantastic and imaginary. By the mid-2010s, however, “hallucinate” begins to acquire negative connotations, being increasingly applied to factually inaccurate output in both computer vision systems and LLMs (Maleki et al., 2024). Hicks et al. unfortunately do not account for nor address this lineage. Consequently, their rather quick dismissal of the term as simply being “the wrong metaphor” (Hicks et al., 2024, p. 9) is a judgment that proceeds on the basis of an incomplete inventory of the rich history of the metaphor with the field of AI. If “hallucinate” is for Hicks et al. “too anthropomorphic,” it is perhaps because their restricted understanding of the term is already too anthropomorphic.

Confabulation. Humans confabulate when, without intending to deceive, they confidently use invented information, in order to, among other things, plug memory lapses in the process of giving information or narrating a sequence of events (Hirstein, 2009). This kind of activity might make some metaphorical sense for LLMs, such as when for example an LLM presents a true statement of fact but, when questioned about it, provides plausible-sounding fictional reasons for that fact. Like “hallucination,” this word is also a metaphorical borrowing from human experience, but unlike “hallucination” it has been identified as more attentive, or at least less problematic, and for the following three reasons:

First, confabulation does not pertain to precepts or stimuli, thus avoiding the thorny implication that LLMs have sensory experiences or even consciousness. Second, as a more neutral term, confabulation is more sensitive to groups dealing with mental health

disorders and better at signaling a lack of malicious intent in that the output is not to be thought of as a fabrication. Finally, the concept’s applicability to non-pathological scenarios is supported by research from psychiatry suggesting that everyday memory reconstruction often involves some degree of confabulation and is equally common for healthy individuals to inadvertently fictionalize details of stories without the intention to deceive (Sui et al., 2024).

We have argued that anthropomorphic language can be reasonable and perhaps informative even if it is metaphorical or analogical. Hicks et al. appear to implicitly recognize this fact. Indeed, while proposing that ChatGPT may be a “hard” bullshitter that can “intend to deceive its listener” about its agenda, they confess at one point that their notion of “hard” bullshit for LLMs is “perhaps not literal” (Hicks et al., 2024, p. 7). They also deploy and approve of Dennett’s (1989) well-known heuristic proposal of adopting the “intentional stance” to characterize, explain, and predict the behavior of entities that have no genuine intentions or rationality (Hicks et al., 2024, p. 8). Thus, it seems they agree after all that some (carefully selected) anthropomorphic language is entirely appropriate for describing LLMs.

As we have also argued, such language can get part of its meaning from background or contextual features. Hicks et al. take this route when they argue that ChatGPT can engage in deceptive “hard” bullshit on the grounds that it was *designed* by humans to “mimic human speech” and “give the impression of concern for truth” (Hicks et al., 2024, 2–3). But it is also true that ChatGPT was developed by means of painstaking pretraining and reinforcement learning to generate accurate and sensible answers to questions (even if, once again, it sometimes or often fails) (see Coeckelbergh & Gunkel, 2025, especially ch. 1 “LLM 101”). That fact might contribute to the meaningfulness and usefulness of the terms “fabrication,” “hallucination,” and “confabulation” and perhaps some others, despite the evident differences between LLMs and human beings.

We have already seen the way in which some of Hicks et al.’s own arguments and assumptions do not rule out and can even pave the way for using anthropomorphic terminology and metaphors to describe the activity of LLMs. However, we certainly do not want to deny that anthropomorphic terms, which are in some ways descriptively useful, might also have various kinds of problems. Caution and judiciousness are appropriate (Smith et al., 2023). Some anthropomorphic words and metaphors may simply be inaccurate or make no sense. Others may risk conflating levels of description, leading one to think, for instance, that cognitively colored terms like “hallucination” or “confabulation” refer to certain internal processes of the LLM rather than to its

exhibited behaviors. And some terms might lead people to place too much trust in LLMs or to blame them rather than their creators (Hicks et al., 2024).

As another anthropomorphic term, “bullshit”—at least as it has been characterized by Frankfurt and adopted by Hicks et al.—is not immune to certain objections. There is a risk that in promulgating the idea that LLMs are “bullshit machines”—irredeemable bullshitters through-and-through—we might induce the public to underestimate rather than overestimate the ability of LLMs to produce sound information and other helpful outputs (although some will see that as a good thing or as something worth risking in order to avoid the worse outcome of overestimating LLMs). More importantly, however, there is the risk that exaggerated warnings, which no-one fully believes, or are themselves problematic to begin with, or which appear increasingly out of touch with the technology as LLMs improve, will be educationally and practically counterproductive. Such risks should be factored into the laudable attempt to craft “good science and technology communication” (Hicks et al., 2024, p. 9) about the often misunderstood and overhyped field of AI.

Conclusions and future direction

None of this is to deny that Hicks et al. offer thoughtful and thought-provoking reflections on both bullshit and the difficult but important question of how to best describe LLMs. To avoid any further misunderstanding, we also agree that many of the criticisms of ChatGPT and other LLMs must be taken seriously: These technologies are a form of autocompleting or pattern recognition lacking any consistently reliable connection to truth; their architectures differ from good old fashioned artificial intelligence (GOF AI), which operated with symbols and principles of logic; they can exhibit glaring lapses in commonsense; they have, in contrast to humans and other animals, no sensory apparatus allowing direct access to the world of phenomena; they cannot be trusted to tell the truth about many things; they can produce problematically biased statements; they can be recruited to spread misinformation; and so on. Our analysis, then, is compatible with various strong criticisms of existing LLMs such as ChatGPT, Claude, Gemini, and Llama.

LLMs and similar forms of AI are powerful technologies, and we need to be attentive to the exigencies of their actual operations and critical of their material, social, epistemological, and moral consequences (see, e.g., Crawford, 2021). But just as powerful are the words—the names, terminologies, metaphors, analogies, and descriptions—that we select and employ in these various endeavors. We are fully in agreement with Hicks et al. about the power of

language. As we also hope to have shown, it is no straightforward task to decide what terminology can be truthfully and responsibly used to describe, characterize, and evaluate what ChatGPT and other LLMs do (or do not do). Critical open-mindedness, some historical awareness of common reactions to new technology, and a nuanced understanding of how various words operate in context can all play a part in meeting this goal.

We have attempted to illustrate one way of going about critically evaluating various terminology and methods of understanding AI. In the first section, we provided a brief review of existing literature, demonstrating how a number of researchers and commentators—either before or contemporary with Hicks et al. (2024)—have also sought to marshal Harry Frankfurt’s activity-centered conception of bullshit to describe and critique LLMs like ChatGPT, thereby reflecting something of a *Zeitgeist* response to this technology. In giving that brief review, we sought not only to situate the essay “ChatGPT is Bullshit” in context but derived from that effort a better understanding of the exact contribution Hicks et al.’s article makes to the existing literature on the subject. Then, in Sects. 2 and 3 of this paper, we argued that Hicks and colleagues, along with some other commentators, are too quick to embrace the idea that LLMs are bullshitters and bullshit machines and, relatedly, too quick to dismiss other—equally anthropomorphic and metaphoric—terms such as “hallucination” and “confabulation.” Their arguments, while making thoughtful contributions to the ongoing study of both LLMs and the concept of bullshit, were found to lack a degree of nuance, coherence, and attention.

Given the growing social significance of LLMs and the importance of word choice, what terminology should we ultimately use or defend? Apart from undertaking the kinds of conceptual reflections we have offered in this paper, there also appears to be a role here for empirical studies that investigate the reception and effects of various words and descriptions on both the lay public and the AI-savvy expert. It may be, for example, that the terms which worry Hicks et al. and various other commentators, do not necessarily engender naïve trust or dangerous misunderstandings. For instance, would we really be inclined to overtrust an agent—human or machine—which we believe is prone to hallucination, fabrication, and confabulation when we ask it for accurate information? It seems unlikely, or at least reasonably uncertain. Still, if some terms really do tend to sow confusion and, what would be worse, to threaten socially dangerous outcomes, then it is important to know how and why that is the case. We can then avoid using those terms, or use them more carefully, and we can try to develop a more helpful vocabulary for understanding and describing ChatGPT, other LLMs, and other forms of generative AI.

Acknowledgements No Acknowledgements. No Funding sources to acknowledge.

Author contributions DJG and SC researched and wrote the manuscript.

Funding Open Access funding enabled and organized by CAUL and its Member Institutions

Data availability No datasets were generated or analysed during the current study.

Declarations

Competing interests The authors declare no competing interests.

Open Access This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by/4.0/>.

References

- Allen, K. (2014). Hallucination and imagination. *Australasian Journal of Philosophy*, 93(2), 287–302. <https://doi.org/10.1080/00048402.2014.984312>
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? In Conference on Fairness, Accountability, and Transparency (FAccT '21), March 3–10, 2021, Virtual Event, Canada. *ACM*. <https://doi.org/10.1145/3442188.3445922>
- Coeckelbergh, M., & Gunkel, D. J. (2025). *Communicative AI: A critical introduction to large Language models*. Polity Books.
- Coghlan, S. (2024). Anthropomorphizing machines: Reality or popular Myth? *Minds and Machines*, 34(3), 25. <https://doi.org/10.1007/s11023-024-09686-w>
- Cohen, G. A. (2002). Deeper into bullshit. In S. Buss, & L. Overton (Eds.), *Contours of agency: Essays on themes from Harry Frankfurt* (pp. 321–339). MIT Press.
- Crawford, K. (2021). *Atlas of AI: Power, politics, and the planetary costs of artificial intelligence*. Yale University Press.
- Dennett, D. C. (1989). *The intentional stance*. MIT Press.
- Derrida, J. (1992). *Act of Literature*. Ed. by Derek Attridge. New York: Routledge.
- Edwards, B. (2023). Why ChatGPT and Bing Chat are so Good at Making Things Up. *Ars Technica*. <https://arstechnica.com/information-technology/2023/04/why-ai-chatbots-are-the-ultimate-bs-machines-and-how-people-hope-to-fix-them/>
- Emsley, R. (2023). ChatGPT: These are not Hallucinations—They're Fabrications and Falsifications. *Schizophrenia* 9(52). <https://doi.org/10.1038/s41537-023-00379-4>
- Frankfurt, H. (1986). On bullshit. *Raritan*, 6(2), 81–100. <https://raritanquarterly.rutgers.edu/issue-index/all-volumes-issues/volume-06/volume-06-number-2>
- Frankfurt, H. (2002). Reply to G. A. Cohen. In S. Buss, & L. Overton (Eds.), *Contours of agency: Essays on themes from Harry Frankfurt* (pp. 339–344). MIT Press.
- Frankfurt, H. (2005). *On bullshit*. Princeton University Press.
- Hicks, M., Townsen, J., Humphries, & Slater, J. (2024). ChatGPT is bullshit. *Ethics and Information Technology*, 26(2), 38, 1–10. <https://doi.org/10.1007/s10676-024-09775-5>
- Hirstein, W. (Ed.). (2009). *Confabulation: Views from neuroscience, psychiatry, psychology, and philosophy*. Oxford University Press.
- Johnson, A. (2010). A new take on deceptive advertising: Beyond Frankfurt's analysis of BS. *Business & Professional Ethics Journal*, 29(1/4), 5–32. <http://www.jstor.org/stable/41340837>
- Kirschenbaum, M. (2023). Prepare for the Textpocalypse. *The Atlantic*. 8 March. <https://www.theatlantic.com/technology/archive/2023/03/ai-chatgpt-writing-language-models/673318/>
- Kolassa, S. (2023). ChatGPT as a Generator of Mindless Bullshit. LinkedIn. 14 October. <https://www.linkedin.com/pulse/chatgpt-generator-mindless-bullshit-stephan-kolassa/>
- Lakshmanan, V. (2022). Why Large Language Models like ChatGPT are Bullshit Artists, and How to Use Them Effectively Anyway. LinkedIn. 15 December. <https://www.linkedin.com/pulse/why-large-language-models-like-chatgpt-bullshit-how-use-lakshmanan/>
- Levinstein, B. A., & Herrmann, D. A. (2024). Still no lie detector for Language models: Probing empirical and conceptual roadblocks. *Philosophical Studies*. <https://doi.org/10.1007/s11098-023-02094-3>
- Levy, N. (2023). *Philosophy, bullshit, and peer review: Elements in epistemology*. Cambridge University Press.
- Mahon, J. E. (2016). The Definition of Lying and Deception. In E. N. Zalta (Ed.), *The Stanford Encyclopedia of Philosophy* (Winter 2016 Edition). <https://plato.stanford.edu/archives/win2016/entries/lying-definition/>
- Maleki, N., Padmanabhan, B., & Dutta, K. (2024). AI Hallucinations: A Misnomer Worth Clarifying. Cornell University, ArXiv, Computer Science, January 9. <https://arxiv.org/html/2401.06796v1>
- Marcus, G. (2024). Evidence that LLMs are reaching a point of diminishing returns — and what that might mean. *Marcus on AI* 13 April. <https://garymarcus.substack.com/p/evidence-that-llms-are-reaching-a>
- Marcus, G., & Davis, E. (2020). GPT-3, Bloviator: OpenAI's Language Generator has no Idea What it's Talking About. *MIT Technology Review*. 22 August. <https://www.technologyreview.com/2020/08/22/1007539/gpt3-openai-language-generator-artificial-intelligence-ai-opinion/>
- McQuillan, D. (2023). ChatGPT is a Bullshit Generator Waging Class War. *Vice*. 10 February. <https://www.vice.com/en/article/akex34/chatgpt-is-a-bullshit-generator-waging-class-war>
- Millière, R., & Buckner, C. (2024). A Philosophical Introduction to Language Models-Part II: The Way Forward. *arXiv preprint arXiv:2405.03207*.
- Narayanan, & Arvind and Sayash Kapoor. (2022). ChatGPT is a Bullshit Generator. But it Can Still be Amazingly Useful. *AI Snake Oil*. 6 December. <https://www.aisnakeoil.com/p/chatgpt-is-a-bullshit-generator-but>
- Narayanan, & Arvind and Sayash Kapoor. (2024). *AI snake oil: what artificial intelligence can do, what it can't, and how to tell the difference*. Princeton University Press.
- Pennycook, G., Cheyne, J. A., Barr, N., Koehler, D. J., & Fugelsang, J. A. (2015). On the reception and detection of Pseudo-Profound bullshit. *Judgment and Decision Making*, 10(6), 549–563. <https://doi.org/10.1017/S1930297500006999>
- Pepperberg, I. M. (2012). Symbolic communication in the grey Parrot. In J. Vonk, & T. K. Shackelford (Eds.), *The Oxford handbook*

- of comparative evolutionary psychology (pp. 297–319). Oxford University Press.
- Ricœur, P. (1984). Time and Narrative, vol. 1. Trans. by Kathleen McLaughlin and David Pellauer. Chicago, IL: University of Chicago Press.
- Romero, A. (2023). ChatGPT: A Bullshit Tool For Bullshit Jobs. *The Algorithmic Bridge*. 12 July. <https://www.thealgorithmicbridge.com/p/chatgpt-a-bullshit-tool-for-bullshit>
- Rudolph, J., Tan, S., & Tan, S. (2023). ChatGPT: Bullshit spewer or the end of traditional assessments in higher education?? *Journal of Applied Learning and Teaching*, 6(1), 342–363. <https://doi.org/10.37074/jalt.2023.6.1.9>
- Slater, J., Humphries, J., & Hicks, M. T. (2024, July 17). ChatGPT Isn't 'Hallucinating'—It's Bullshitting! *Scientific American*. <https://www.scientificamerican.com/article/chatgpt-isnt-hallucinating-its-bullshitting/>. Accessed on 10 December 2024.
- Smith, A., Greaves, F., & Trishan Panch. (2023). Hallucination or confabulation?? Neuroanatomy as metaphor in large Language models. *PLOS Digit Health*, 2(11), e0000388. <https://doi.org/10.1371/journal.pdig.0000388>
- Sui, P., Duede, E., Wu, S., & Richard Jean So. (2024). and Confabulation: The Surprising Value of Large Language Model Hallucinations. *Arxiv CS*. <https://arxiv.org/abs/2406.04175>
- Zhang, M., Press, O., Merrill, W., Liu, A., & Smith, N. A. (2023). How Language Model Hallucinations Can Snowball. *arXiv preprint arXiv:2305.13534*.

Publisher's note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.